

LambdaMART notes

Yafei Zhang

August 14, 2014

Contents

1	Notations	2
2	Definitions	3
3	Sum them together	4
4	Gradient Boosting	5
5	Gradient Boosted Regression Trees	5
6	LambdaMART	6

1 Notations

x, x_i, x_j : input feature vector, $x \in \mathbb{R}^n$.

y, y_i, y_j : input score, measurement of goodness.

U, U_i, U_j : URLs.

$s \equiv F(x)$, $s_i \equiv F(x_i)$, $s_j \equiv F(x_j)$: model learned score, measurement of goodness.
 F is the learned model.

$o_{ij} \equiv s_i - s_j$ (for a given query).

$\rho_{ij} \equiv \frac{1}{1+e^{o_{ij}}}$ (for a given query).

P_{ij} : a learned probability that U_i should be ranked higher than U_j (for a given query).

\bar{P}_{ij} : the known probability that U_i should be ranked higher than U_j (for a given query).

S_{ij} : it is defined as (for a given query)

$$S_{ij} \equiv \begin{cases} 0 & \text{for } y_i = y_j \\ 1 & \text{for } y_i > y_j \\ -1 & \text{for } y_i < y_j \end{cases} \quad (1)$$

C : target function we want to maximize.

C_{ij} : the amount that U_i and U_j contribute to C (for a given query).

λ_{ij} : lambda gradient of U_i and U_j (for a given query).

λ_i : sums of lambda gradient of U_i and all other URLs.

$I \equiv \{(i, j) | S_{ij} = 1\}$: the set of pairs of indices (i, j) , for which we desire U_i to be ranked higher than U_j (for a given query).

Z : a measure of L2R such as NDCG, MAP, MRR. It is NDCG in LambdaMART.

ΔZ_{ij} is the change of Z by swapping the rank positions of U_i and U_j (for a given query) (after sorting all documents by their current scores s).

2 Definitions

$$P_{ij} \equiv P(U_i \succ U_j) \equiv \frac{1}{1 + e^{-(s_i - s_j)}} = \frac{1}{1 + e^{-o_{ij}}} \quad (2)$$

$$\bar{P}_{ij} \equiv \frac{(1 + S_{ij})}{2} \quad (3)$$

C_{ij} is chosen to be:

$$C_{ij} \equiv |\Delta Z_{ij}| \left[-\bar{P}_{ij} \log(P_{ij}) - (1 - \bar{P}_{ij}) \log(1 - P_{ij}) \right] \quad (4)$$

$$= |\Delta Z_{ij}| \left[\frac{1}{2} (1 - S_{ij}) o_{ij} + \log(1 + e^{-o_{ij}}) \right] \quad (5)$$

$$= \begin{cases} |\Delta Z_{ij}| \log(1 + e^{-o_{ij}}) & \text{for } S_{ij} = 1 \\ |\Delta Z_{ij}| \log(1 + e^{-o_{ji}}) & \text{for } S_{ij} = -1 \end{cases} \quad (6)$$

From now on, we assume all $(i, j) \in I$ except extra explanations.

$$C_{ij} \equiv |\Delta Z_{ij}| \log(1 + e^{-o_{ij}}) \quad (7)$$

Gradient(treat ΔZ_{ij} a constant):

$$\frac{\partial C_{ij}}{\partial s_i} = \frac{\partial C_{ij}}{\partial o_{ij}} = -\frac{|\Delta Z_{ij}|}{1 + e^{o_{ij}}} = -|\Delta Z_{ij}| \rho_{ij} = -\frac{\partial C_{ij}}{\partial o_{ji}} = -\frac{\partial C_{ij}}{\partial s_j} \quad (8)$$

λ -gradient:

$$\lambda_{ij} \equiv \left| \frac{\partial C_{ij}}{\partial o_{ij}} \right| = \frac{|\Delta Z_{ij}|}{1 + e^{o_{ij}}} = |\Delta Z_{ij}| \rho_{ij} = -\lambda_{ji} \quad (9)$$

So

$$\frac{\partial C_{ij}}{\partial s_i} = -\lambda_{ij} \quad (10)$$

$$\frac{\partial^2 C_{ij}}{\partial s_i^2} = -\frac{\partial \lambda_{ij}}{\partial s_i} = \frac{|\Delta Z| e^{o_{ij}}}{(1 + e^{o_{ij}})^2} = |\Delta Z| \rho_{ij} (1 - \rho_{ij}) \quad (11)$$

3 Sum them together

$$C = \sum_i \sum_{j:(i,j) \in I} C_{ij} + \sum_i \sum_{j:(j,i) \in I} C_{ji} \quad (12)$$

$$= \sum_i \sum_{j:(i,j) \in I} |\Delta Z_{ij}| \log(1 + e^{-o_{ij}}) + \sum_i \sum_{j:(j,i) \in I} |\Delta Z_{ji}| \log(1 + e^{-o_{ji}}) \quad (13)$$

$$\frac{\partial C}{\partial s_i} = \sum_{j:(i,j) \in I} \frac{\partial C_{ij}}{\partial s_i} + \sum_{j:(j,i) \in I} \frac{\partial C_{ji}}{\partial s_i} \quad (14)$$

$$= \sum_{j:(i,j) \in I} (-|\Delta Z_{ij}| \rho_{ij}) + \sum_{j:(j,i) \in I} |\Delta Z_{ji}| \rho_{ji} \quad (15)$$

$$= \sum_{j:(i,j) \in I} (-\lambda_{ij}) + \sum_{j:(j,i) \in I} \lambda_{ji} \quad (16)$$

$$\lambda_i \equiv \sum_{j:(i,j) \in I} \lambda_{ij} + \sum_{j:(j,i) \in I} \lambda_{ij} \quad (17)$$

$$= \sum_{j:(i,j) \in I} \lambda_{ij} - \sum_{j:(j,i) \in I} \lambda_{ji} \quad (18)$$

So

$$\frac{\partial C}{\partial s_i} = -\lambda_i \quad (19)$$

$$\frac{\partial^2 C}{\partial s_i^2} = \sum_{j:(i,j) \in I} \frac{\partial^2 C_{ij}}{\partial s_i^2} + \sum_{j:(j,i) \in I} \frac{\partial^2 C_{ji}}{\partial s_i^2} \quad (20)$$

$$= \sum_{j:(i,j) \in I} |\Delta Z_{ij}| \rho_{ij} (1 - \rho_{ij}) + \sum_{j:(j,i) \in I} |\Delta Z_{ji}| \rho_{ji} (1 - \rho_{ji}) \quad (21)$$

$$= \sum_{j:(i,j) \in I} \left(-\frac{\partial \lambda_{ij}}{\partial s_i}\right) + \sum_{j:(j,i) \in I} \left(-\frac{\partial \lambda_{ji}}{\partial s_j}\right) \quad (22)$$

$$= \sum_{j:(i,j) \in I} \left(-\frac{\partial \lambda_{ij}}{\partial s_i}\right) + \sum_{j:(j,i) \in I} \frac{\partial \lambda_{ji}}{\partial s_i} \quad (23)$$

So

$$\frac{\partial^2 C}{\partial s_i^2} = -\frac{\partial \lambda_i}{\partial s_i} \quad (24)$$

4 Gradient Boosting

In this and next section, we will take a digression to recall some knowledge of GBRT(Gradient Boosted Regression Trees).

We consider L a function of input score y and model function $F = F(x)$: $L = L(y, F)$. Our goal is

$$\min_F L(y, F) \quad (25)$$

As in ordinary gradient descent, F is iteratively updated in functional space.

$$F_{n+1} = F_n - \rho \left. \frac{\partial L(y, F)}{\partial F} \right|_{F=F_n} \quad (26)$$

ρ is a step length, chosen

$$\rho = \operatorname{argmin}_\rho C\left(y, F_n - \rho \left. \frac{\partial L(y, F)}{\partial F} \right|_{F=F_n}\right) \quad (27)$$

26 and 27 can be re-interpreted as

$$F_{n+1} = F_n + \rho f_n(x) \quad (28)$$

where $f_n(x)$ is a model that fits pseudo-response $\{\widetilde{y}\}_i$.

$$\widetilde{y}_i = - \left. \frac{\partial L(y_i, F)}{\partial F} \right|_{F=F_n} \quad (29)$$

Then ρ is chosen

$$\rho = \operatorname{argmin}_\rho L(y, F_n + \rho f_n(x)) \quad (30)$$

5 Gradient Boosted Regression Trees

Particularly, if we choose $f(x)$ as a regression tree, this results in a GBRT algorithm.

$$f(x) = f(x, \{\gamma_j, R_j\}_i^J) = \sum_{j=1}^J \gamma_j 1(x \in R_j) \quad (31)$$

J is the number of leaves. $\{R_j\}_i^J$ are disjoint regions that cover all feature space, each of them is covered in one leaf. $\{\gamma_j\}_i^J$ are the values in the corresponding leaf.

In this circumstance, we first construct $\{R_j\}_i^J$ to fit $\{x_i, \widetilde{y}\}_i$ by least-squares, then use 26,

$$F_{n+1} = F_n + \sum_{j=1}^J \gamma_j 1(x \in R_j) \quad (32)$$

use 27.

$$\gamma_j = \operatorname{argmin}_{\gamma} \sum_{i: x_i \in R_j} L(y_i, F_n(x_i) + \gamma) = \operatorname{argmin}_{\gamma} \sum_{i: x_i \in R_j} g \quad (33)$$

where

$$g = g(y_i, F_n(x_i)) = L(y_i, F_n(x_i) + \gamma) \quad (34)$$

When there is no closed form solution to 33, we can approximate it by a single Newton step. This is

$$\gamma_j = - \frac{\sum_{i: x_i \in R_j} \frac{\partial g}{\partial F_n}}{\sum_{i: x_i \in R_j} \frac{\partial^2 g}{\partial F_n^2}} \quad (35)$$

GBRT is sometimes called GBDT(Gradient Boosted Decision Trees) or MART(Multi Additive Regression Trees).

Note that GBDT is so-called for its decision purpose, but each tree in it is still a regression tree.

6 LambdaMART

Our goal is

$$\min_F (-C) \quad (36)$$

with $-C$ defined in 13, with the first order derivative $-\frac{\partial C}{\partial s_i}$ defined in 19, with the second order derivative $-\frac{\partial^2 C}{\partial s_i^2}$ defined in 24.

F is MART. Overall training framework is depicted below.

```

input : Training samples  $\{x_i, y_i, U_i\}_1^N$  ( $x_i$  contains a query), the metric  $Z$ , the number of
        leaves in a tree  $L$ , the number of trees  $M$ , learning rate  $\nu$ 
output:  $F, \{s_i\}_1^N$ 
data : pseudo responses  $\{\tilde{y}_i\}_1^N$ , weights  $\{w_i\}_1^N$ , model output  $\{s_i\}_1^N$ ;
1 for  $i = 1$  to  $N$  do
2    $s_i = F_0(x_i) = \text{BaseModel}(x_i)$ ;
3 end
4 for  $m = 1$  to  $M$  do
5   for  $i = 1$  to  $N$  do
6      $\tilde{y}_i = \lambda_i$  defined in 19;
7      $w_i = \frac{\partial \lambda_i}{\partial s_i}$  defined in 24;
8   end
9   Create a tree  $\{R_{lm}\}_{l=1}^L$  to fit  $\{x_i, \tilde{y}_i\}_1^N$  by least-squares;
10  for  $l = 1$  to  $L$  do
11     $\gamma_{lm} = -\frac{\sum_{i: x_i \in R_{lm}} \tilde{y}_i}{\sum_{i: x_i \in R_{lm}} w_i}$ ;
12  end
13  for  $i = 1$  to  $N$  do
14     $s_i = F_m(x_i) = F_{m-1}(x_i) + \nu \sum_l \gamma_{lm} 1(x_i \in R_{lm})$ ;
15  end
16 end
17  $F = F_M$ ;

```

Algorithm 1: The LambdaMART algorithm

References

- [1] Qiang Wu, Christopher J.C. Burges, Krysta M. Svore, Jianfeng Gao. Adapting Boosting for Information Retrieval Measures. Learning to Rank for Information Retrieval DOI 10.1007/s10791-009-9112-1, 2009
- [2] Christopher J.C. Burges. From RankNet to LambdaRank to LambdaMART: An Overview. Microsoft Research Technical Report MSR-TR-2010-82, 2010.
- [3] Jerome H. Friedman. Greedy function Approximation: A Gradient Boosting Machine. Annals of Statistics, Vol. 29 (2001), pp. 1189-1232, 1999.